


RAG Demo - Architecture & Implementation Guide

System Architecture

High-Level Overview



Components

1. Frontend (Angular 18)

- **Chat Interface:** User-facing Q&A interface
- **Admin Panel:** Document upload and website scraping
- **Services:** HTTP communication with backend API
- **Features:**
 - Real-time chat with loading states
 - Markdown formatting (lists, headings, bold, code)
 - Conversational message handling
 - Document upload (PDF)
 - Website content ingestion

2. Backend (ASP.NET Core 8)

- **Controllers:** REST API endpoints
- **Services:**
 - Chat Service: AI response generation
 - Document Service: PDF and web content processing
 - Embedding Service: Text vectorization
 - Vector Store Service: Qdrant integration
- **Features:**
 - RAG (Retrieval-Augmented Generation)
 - Vector similarity search
 - PDF text extraction
 - Website scraping with depth control
 - Conversational context detection

3. Vector Database (Qdrant)

- Stores document embeddings (768-dimensional vectors)
- Semantic similarity search
- Metadata storage for source tracking

4. AI Model (GitHub Models / Azure AI)

- Model: GPT-4o-mini
 - Used for intelligent response generation
 - Fallback to mock responses if unavailable
-

Implementation Steps

Phase 1: Core RAG Pipeline

1. Document Ingestion

- PDF text extraction using iTextSharp
- Text chunking (configurable size with overlap)
- Generate embeddings using local model

2. Vector Storage

- Store chunks in Qdrant with metadata
- Index by document source and content

3. Retrieval & Generation

- User asks a question
- Generate query embedding
- Search top-K similar chunks
- Pass context to LLM
- Generate answer with sources

Phase 2: Frontend Development

1. Chat Interface

- Message display with user/bot distinction
- Markdown rendering with icons
- Loading states and error handling
- Change detection optimization (OnPush strategy)

2. Admin Panel

- PDF document upload
- Website URL ingestion
- Progress indicators
- Success/error notifications

3. Routing

- Separate routes for chat (/chat) and admin (/admin)
- Navigation bar

Phase 3: Enhancements

1. Smart Conversational Detection

- Filter casual greetings ("thanks", "hi", "bye")
- Avoid unnecessary document searches

2. Response Formatting

- Structured prompts for better AI output
- Headings, lists, bold text support
- Code snippet formatting

3. Website Scraping

- Extract text from web pages
- Optional link crawling (1-3 levels)
- Multiple page processing

Key Technologies

Component	Technology	Purpose
Frontend	Angular 18 (Standalone)	Modern SPA framework
UI Library	PrimeNG	Professional UI components
Backend	ASP.NET Core 8	RESTful API
Vector DB	Qdrant	Semantic search
Embeddings	all-MiniLM-L6-v2	Local embedding model
AI Model	GPT-4o-mini (GitHub)	Response generation
PDF Parser	iTextSharp	Text extraction
Web Scraper	HtmlAgilityPack	HTML parsing

Production Readiness Checklist

1. Security

- Add authentication/authorization (JWT tokens)
- Implement rate limiting on API endpoints
- Add CORS configuration for specific domains
- Secure API keys in Azure Key Vault or similar

- Input validation and sanitization
- HTTPS enforcement
- SQL injection protection (use parameterized queries)

2. Performance

- Enable response caching for repeated queries
- Implement connection pooling for Qdrant
- Add CDN for frontend static assets
- Optimize embedding batch processing
- Implement request throttling
- Database indexing optimization
- Lazy loading for frontend modules

3. Monitoring & Logging

- Application Insights / ELK Stack integration
- API endpoint monitoring
- Error tracking (Sentry or similar)
- Performance metrics (response times, throughput)
- Vector store health checks
- User activity analytics
- Alert system for failures

4. Scalability

- Containerize with Docker
- Kubernetes orchestration
- Horizontal scaling for API
- Load balancer configuration
- Qdrant cluster setup for high availability
- Message queue for async processing (RabbitMQ)
- Redis for session management

5. Data Management

- Backup strategy for vector database
- Document versioning system
- Data retention policies
- GDPR compliance (data deletion)
- Audit logs for document operations

6. Quality Assurance

- Unit tests (>80% coverage)
- Integration tests for RAG pipeline
- E2E tests for critical user flows
- Load testing (Apache JMeter / k6)
- Security testing (OWASP)

- API documentation (Swagger/OpenAPI)

7. Configuration

- Environment-specific configs (Dev/Staging/Prod)
 - Feature flags for gradual rollouts
 - Configurable chunk sizes and overlap
 - Adjustable similarity thresholds
 - Model selection configuration
-

Deployment Guide

Option 1: Azure Deployment (Recommended)

Backend (.NET API)

```
# 1. Create Azure App Service
az webapp create --name rag-demo-api --resource-group rg-rag-demo --plan asp-rag-
demo

# 2. Configure app settings
az webapp config appsettings set --name rag-demo-api --resource-group rg-rag-demo
--settings \
  "GitHub__Token=<your-token>" \
  "Qdrant__Url=<qdrant-url>" \
  "Qdrant__ApiKey=<qdrant-key>"

# 3. Deploy from GitHub Actions
# (See deployment pipeline below)
```

Frontend (Angular)

```
# 1. Build production bundle
npm run build -- --configuration production

# 2. Deploy to Azure Static Web Apps
az staticwebapp create --name rag-demo-ui --resource-group rg-rag-demo

# 3. Configure API URL
# Update environment.production.ts with API endpoint
```

Qdrant Vector Database

```
# Option A: Qdrant Cloud (Managed)
# Sign up at cloud.qdrant.io and create cluster
```

```
# Option B: Self-hosted on Azure Container Instances
az container create --name qdrant --resource-group rg-rag-demo \
--image qdrant/qdrant:latest --ports 6333 6334 \
--cpu 2 --memory 4
```

Option 2: Docker Deployment

```
# docker-compose.yml
version: '3.8'
services:
  frontend:
    build: ./RAGDemoFrontend
    ports:
      - "80:80"
    environment:
      - API_URL=http://backend:5000

  backend:
    build: ./RAGDemoBackend
    ports:
      - "5000:8080"
    environment:
      - Qdrant__Url=http://qdrant:6333
      - GitHub__Token=${GITHUB_TOKEN}
    depends_on:
      - qdrant

  qdrant:
    image: qdrant/qdrant:latest
    ports:
      - "6333:6333"
    volumes:
      - qdrant-data:/qdrant/storage

volumes:
  qdrant-data:
```

```
# Deploy
docker-compose up -d

# Scale backend
docker-compose up -d --scale backend=3
```

Option 3: Kubernetes Deployment

```
# k8s-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: rag-demo-api
spec:
  replicas: 3
  selector:
    matchLabels:
      app: rag-demo-api
  template:
    metadata:
      labels:
        app: rag-demo-api
    spec:
      containers:
        - name: api
          image: your-registry/rag-demo-api:latest
          ports:
            - containerPort: 8080
          env:
            - name: Qdrant__Url
              value: "http://qdrant-service:6333"
          resources:
            requests:
              memory: "512Mi"
              cpu: "500m"
            limits:
              memory: "1Gi"
              cpu: "1000m"
      ...
apiVersion: v1
kind: Service
metadata:
  name: rag-demo-api-service
spec:
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: rag-demo-api
```

CI/CD Pipeline (GitHub Actions)

```
# .github/workflows/deploy.yml
name: Deploy RAG Demo

on:
```

```
push:
  branches: [ main ]

jobs:
  build-and-deploy-backend:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3

      - name: Setup .NET
        uses: actions/setup-dotnet@v3
        with:
          dotnet-version: '8.0.x'

      - name: Build
        run: dotnet build --configuration Release
        working-directory: ./RAGDemoBackend

      - name: Publish
        run: dotnet publish -c Release -o ./publish
        working-directory: ./RAGDemoBackend

      - name: Deploy to Azure
        uses: azure/webapps-deploy@v2
        with:
          app-name: 'rag-demo-api'
          publish-profile: ${{ secrets.AZURE_WEBAPP_PUBLISH_PROFILE }}
          package: './RAGDemoBackend/publish'

build-and-deploy-frontend:
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v3

    - name: Setup Node
      uses: actions/setup-node@v3
      with:
        node-version: '20'

    - name: Install dependencies
      run: npm ci
      working-directory: ./RAGDemoFrontend

    - name: Build
      run: npm run build -- --configuration production
      working-directory: ./RAGDemoFrontend

    - name: Deploy to Azure Static Web Apps
      uses: Azure/static-web-apps-deploy@v1
      with:
        azure_static_web_apps_api_token: ${{ secrets.AZURE_STATIC_WEB_APPS_TOKEN
}}        repo_token: ${{ secrets.GITHUB_TOKEN }}
```

```
action: "upload"
app_location: "./RAGDemoFrontend/dist"
```

Cost Estimation (Azure - Monthly)

Service	Tier	Estimated Cost
App Service (API)	B1 Basic	\$13
Static Web App (Frontend)	Free	\$0
Container Instance (Qdrant)	2 vCPU, 4GB	\$50
GitHub Models API	Pay-per-use	\$5-20
Application Insights	Basic	\$2
Total		\$70-85/month

Cost Optimization Tips

- Use Azure Reserved Instances (save 30-40%)
- Implement caching to reduce AI API calls
- Use Qdrant Cloud free tier for development
- Auto-scaling based on traffic patterns

Success Metrics

Technical KPIs

- Response time < 2 seconds (p95)
- API availability > 99.9%
- Search accuracy > 85%
- Vector search latency < 100ms

Business KPIs

- User engagement rate
- Query success rate
- Document processing throughput
- Cost per query

Next Steps / Roadmap

1. Short Term (1-2 weeks)

- Add user authentication
- Implement request logging

- Create comprehensive API documentation

2. Medium Term (1-2 months)

- Multi-language support
- Advanced search filters
- Document management dashboard
- Chat history persistence

3. Long Term (3-6 months)

- Multi-tenant support
 - Custom embedding models
 - Real-time collaboration
 - Mobile app (React Native)
 - Advanced analytics dashboard
-

Support & Resources

- **Documentation:** See README.md for setup instructions
 - **API Docs:** <https://your-api.com/swagger>
 - **Qdrant Docs:** <https://qdrant.tech/documentation/>
 - **GitHub Models:** <https://github.com/marketplace/models>
 - **PrimeNG:** <https://primeng.org/>
-

Last Updated: January 2026

Version: 1.0

Maintainer: Mohd Waseem