

# RAG Demo - Full-Stack AI-Powered Q&A System

---

A production-ready Retrieval-Augmented Generation (RAG) system with Angular frontend and .NET Core backend, featuring document upload, website scraping, and intelligent conversational AI.

---

## Features Implemented

### Backend Features (.NET 8.0)

- Local Vector Database** - Qdrant for semantic search
- Local Embeddings** - No API costs, runs offline
- PDF Processing** - Extract and chunk documents with metadata
- Website Scraping** - Ingest content from web pages with depth control
- Semantic Search** - Find relevant context using vector similarity
- GitHub Models Integration** - FREE AI-powered responses (GPT-4o-mini)
- Smart Conversational Detection** - Filters casual greetings to avoid unnecessary searches
- Formatted Prompts** - Structured AI responses with lists, headings, and code
- Multi-page Crawling** - Optional link following (1-3 levels deep)
- RESTful API** - Clean, documented endpoints with Swagger
- Cost Effective** - 100% free for local deployment
- Production Ready** - Scalable architecture with proper DI

### Frontend Features (Angular 18)

- Modern Chat Interface** - Real-time Q&A with loading states
  - Markdown Rendering** - Support for lists, headings, bold, code snippets
  - Admin Panel** - Separate interface for document management
  - Document Upload** - PDF file upload with progress indicators
  - Website Ingestion** - URL-based content scraping from admin panel
  - Message Formatting** - Beautiful list icons and structured responses
  - Conversational UX** - Smart handling of greetings and casual messages
  - Routing** - Separate routes for chat (/chat) and admin (/admin)
  - Error Handling** - Toast notifications for success/error states
  - Responsive Design** - Works on desktop, tablet, and mobile
  - PrimeNG UI** - Professional, modern UI components
  - Change Detection Optimization** - OnPush strategy for performance
- 

## Prerequisites

### Required

- .NET 8.0 SDK
- Docker Desktop (for Qdrant)
- Node.js 20+ and npm (for frontend)

## Optional (for AI-powered responses)

- GitHub Models token (FREE) - Get from <https://github.com/marketplace/models>

## 🛠 Quick Start

### 1. Start Qdrant Vector Database

```
# Pull and run Qdrant using Docker
docker run -d -p 6333:6333 -p 6334:6334 ` 
  -v qdrant_storage:/qdrant/storage ` 
  --name qdrant ` 
  qdrant/qdrant
```

Verify Qdrant is running:

```
curl http://localhost:6333/dashboard
```

### 2. Start Backend API

```
cd RAGDemoBackend
dotnet restore
dotnet run
```

The API will start at <https://localhost:5001> or <http://localhost:5000>

### 3. Start Frontend

```
cd RAGDemoFrontend
npm install
npm start
```

The app will open at <http://localhost:4200>

### 4. (Optional) Enable AI-Powered Responses

Get a FREE token from <https://github.com/marketplace/models>

```
# Set environment variable
$env:GH_TOKEN = "your-github-token"

# Or edit appsettings.json
```

```
# "GitHub": { "Token": "your-token" }
# "DemoSettings": { "UseGitHubModels": true }
```

## 5. Use the Application

1. Navigate to **Admin Panel** (<http://localhost:4200/admin>)
2. Upload a PDF document or ingest a website URL
3. Go to **Chat** (<http://localhost:4200/chat>)
4. Ask questions about your documents!

## Architecture



Qdrant Database  
(Docker)  
- Vector Storage  
- Cosine Search

GitHub Models  
(GPT-4o-mini)  
- AI Responses  
- Free Tier

## 🔧 Configuration

Edit `appsettings.json`:

```
{  
    "Qdrant": {  
        "Host": "localhost",  
        "Port": "6334",  
        "UseHttps": false,  
        "CollectionName": "documents"  
    },  
    "DemoSettings": {  
        "ChunkSize": 500,  
        "MaxSearchResults": 5  
    }  
}
```

## 📡 API Endpoints

### Health Check

```
GET /api/chat/health  
Response: {  
    "status": "RAG Demo API is running",  
    "timestamp": "2026-01-15T10:30:00Z"  
}
```

### Get Statistics

```
GET /api/chat/stats  
Response: {  
    "documentChunks": 42,  
    "vectorStore": "Qdrant",  
    "embeddingModel": "Local (all-MiniLM-L6-v2 compatible)",  
    "timestamp": "2026-01-15T10:30:00Z"  
}
```

## Upload Document

```
POST /api/chat/upload
Content-Type: multipart/form-data

file: [PDF file]

Response: {
  "message": "Processed document.pdf"
}
```

## Ingest Website Content

```
POST /api/chat/ingest-url
Content-Type: application/json

{
  "url": "https://example.com",
  "includeLinks": true,
  "maxDepth": 2
}

Response: {
  "url": "https://example.com",
  "chunksCreated": 25,
  "status": "Success",
  "processedUrls": ["https://example.com", "https://example.com/about"]
}
```

## Ask Question

```
POST /api/chat/ask
Content-Type: application/json

{
  "question": "What is the product pricing?",
  "sessionId": "optional-session-id"
}

Response: {
  "answer": "Based on the documents, here are the key pricing details:\n\n1. Basic Plan - $29/month\n2. Pro Plan - $79/month\n3. Enterprise - Custom pricing",
  "sources": ["product-guide.pdf", "pricing.pdf"],
  "sessionId": "abc-123"
}
```

## Delete Document

```
DELETE /api/chat/document/{documentName}
```

```
Response: {
    "message": "Deleted document.pdf"
}
```

## Test with Swagger

Navigate to: <https://localhost:5001/swagger>

## ⌚ How It Works

### 1. Document Ingestion

#### **PDF Upload:**

- PDF uploaded via frontend admin panel
- Text extracted using iText7
- Text split into configurable chunks (default 500 chars)
- Each chunk generates a 768-dimensional embedding vector
- Chunks + embeddings + metadata stored in Qdrant

#### **Website Scraping:**

- URL provided via admin panel
- HTML content fetched and parsed
- Text extracted from main content areas
- Optional: Follow links up to 3 levels deep
- Each page chunked and embedded separately
- Source URL tracked in metadata

### 2. Semantic Search

- User asks a question in chat interface
- Question converted to embedding vector (same model)
- Qdrant performs cosine similarity search
- Top 5 most relevant chunks retrieved
- Source documents tracked for citation

### 3. Intelligent Response Generation

#### **Conversational Detection:**

- System filters casual messages ("thanks", "hi", "bye")
- Avoids unnecessary vector searches for greetings
- Provides friendly conversational responses

## RAG Pipeline:

- Retrieved chunks provide context to AI
- GitHub Models (GPT-4o-mini) generates response
- Response formatted with markdown (lists, headings, code)
- Sources included for transparency
- Fallback to mock response if AI unavailable

## 4. Frontend Rendering

- Markdown automatically parsed and styled
  - Lists display with icons (✓ for numbered, • for bullets)
  - Headings properly sized and weighted
  - Code snippets highlighted
  - Real-time loading indicators
  - Toast notifications for actions
- 

## ⌚ Cost Analysis

| Component       | Cost             | Notes               |
|-----------------|------------------|---------------------|
| Qdrant (Docker) | \$0              | Runs locally        |
| Embeddings      | \$0              | Local CPU inference |
| PDF Processing  | \$0              | iText7 open source  |
| .NET Hosting    | \$0              | Local dev           |
| <b>Total</b>    | <b>\$0/month</b> | 🚀 Free!             |

## Production Costs (Optional)

| Upgrade              | Monthly Cost | When Needed       |
|----------------------|--------------|-------------------|
| OpenAI API (GPT-3.5) | ~\$5-20      | Better answers    |
| Azure App Service    | ~\$55        | Cloud hosting     |
| Qdrant Cloud (1GB)   | ~\$25        | Managed vector DB |

---

## 🚀 Key Features Breakdown

### Chat Interface

- **Real-time messaging** with user/bot distinction
- **Markdown support** for formatted responses
- **List rendering** with visual icons (✓ • )
- **Code highlighting** for technical content
- **Loading states** with spinner animation
- **Error handling** with user-friendly messages

- **Keyboard shortcuts** (Enter to send, Shift+Enter for new line)
- **Auto-scroll** to latest message
- **Session persistence** across questions

## Admin Panel

- **Dual upload methods:** PDF files or website URLs
- **Progress tracking** for uploads and ingestion
- **Depth control** for website crawling (1-3 levels)
- **Link following** option for comprehensive scraping
- **Success notifications** with chunk counts
- **Processed URL list** showing all scraped pages
- **Responsive design** for mobile admin access

## Backend Intelligence

- **Smart routing** - Conversational vs. knowledge queries
  - **Contextual prompts** - Structured AI instructions
  - **Source tracking** - Maintains document provenance
  - **Fallback handling** - Mock responses if AI unavailable
  - **Chunk optimization** - Configurable size and overlap
  - **Metadata enrichment** - URL, title, timestamp tracking
- 

## 💡 Next Steps & Roadmap

### Current Capabilities

- Document Q&A with PDF uploads
- Website content ingestion
- AI-powered responses with formatting
- Separate admin and user interfaces
- Real-time chat experience
- Source attribution

### Immediate Improvements

- Add user authentication (JWT)
- Implement chat history persistence
- Add document management dashboard
- Support more file formats (DOCX, TXT)
- Add batch upload capability
- Implement response streaming

### Production Enhancements

- Database persistence (Entity Framework + PostgreSQL)
- Redis caching for frequent queries
- Rate limiting per user/IP

- Multi-tenant support
  - Advanced analytics dashboard
  - Multi-language support
  - Mobile app (React Native)
  - Document versioning
  - Custom embedding models
  - A/B testing for prompts
- 

## ❖ Tips & Best Practices

### For Best Results

1. **Chunk Size:** 500 chars works well for general content. Increase to 1000 for technical documentation.
2. **Search Results:** 5 chunks provide good context. Adjust in `appsettings.json` if needed.
3. **Question Format:** Be specific in questions for better retrieval accuracy.
4. **Document Organization:** Upload related documents together for coherent context.
5. **Website Scraping:** Start with `maxDepth: 1` to test before going deeper.
6. **Memory:** Qdrant caches chunks in memory for fast retrieval.

### Performance Optimization

- Use `ChangeDetectionStrategy.OnPush` in Angular components
- Implement lazy loading for large document lists
- Cache frequent queries on backend
- Optimize chunk size based on content type
- Use indexes on Qdrant collections

### Security Considerations

- Validate file uploads (size, type, content)
  - Sanitize user inputs before querying
  - Implement rate limiting on API endpoints
  - Use HTTPS in production
  - Secure API keys in environment variables
  - Add authentication before production deployment
- 

## 📖 Documentation

- **ARCHITECTURE.md** - Detailed system architecture and production deployment guide
  - **API Documentation** - Available at <https://localhost:5001/swagger> when backend is running
  - **Component Structure** - See `/src/app/components` for Angular component organization
- 

## 🤝 Contributing

This is a demo/reference project showcasing RAG implementation. For production use:

- Add comprehensive unit and integration tests

- Implement proper error handling and logging
  - Add monitoring and alerting
  - Security hardening (authentication, authorization, input validation)
  - Performance optimization and load testing
  - Documentation for API consumers
- 

## License

MIT License - Feel free to use in your projects!

---

## Learn More

- **RAG Concepts:** <https://www.pinecone.io/learn/retrieval-augmented-generation/>
  - **Qdrant Docs:** <https://qdrant.tech/documentation/>
  - **GitHub Models:** <https://github.com/marketplace/models>
  - **Angular Best Practices:** <https://angular.dev/best-practices>
  - **PrimeNG Components:** <https://primeng.org/>
- 

**Questions or Issues?** Check the troubleshooting section or review the logs!

**Production Ready?** See [ARCHITECTURE.md](#) for deployment checklist.

---

Built with  using .NET, Angular, and Qdrant